


Choose an Azure compute service for your application

7/24/2020 • 6 minutes to read •  +4

Steps article

[Choose a candidate service](#)

[Understand the basic features](#)

[Understand the hosting models](#)

[Understand the limits and cost](#)

[Understand the basic features](#)

[Understand the limits and cost](#)

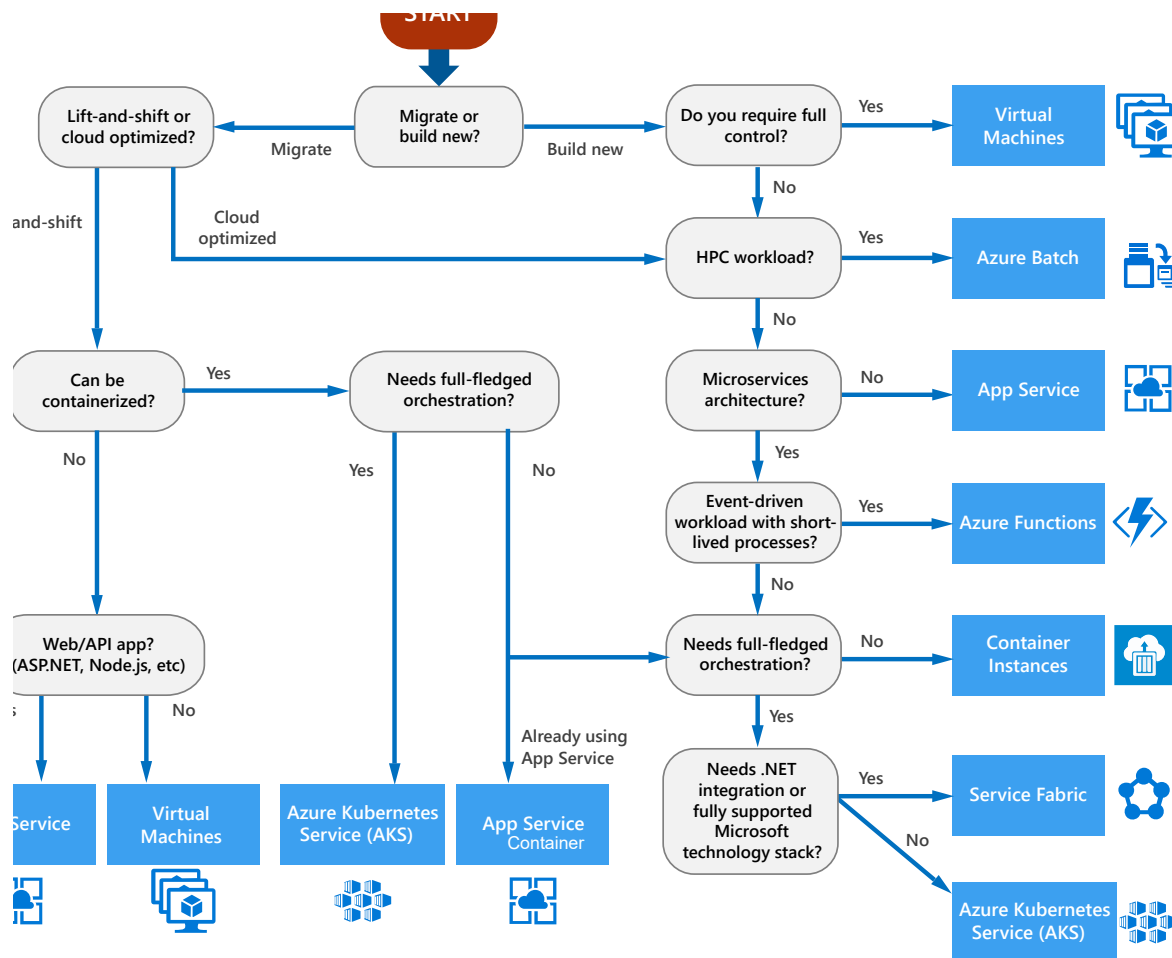
[Next steps](#)

Azure offers a number of ways to host your application code. The term *compute* refers to a hosting model for the computing resources that your application runs on. The following flowchart will help you to choose a compute service for your application.

If your application consists of multiple workloads, evaluate each workload separately. A complete solution may incorporate two or more compute services.

Choose a candidate service

Use the following flowchart to select a candidate compute service.



itions:

"Lift and shift" is a strategy for migrating a workload to the cloud without redesigning the application or making code changes. Also called *rehosting*. For more information, see [Azure migration center](#).

Cloud optimized is a strategy for migrating to the cloud by refactoring an application to take advantage of cloud-native features and capabilities.

Output from this flowchart is a **starting point** for consideration. Next, perform a detailed evaluation of the service to see if it meets your needs.

Understand the basic features

If you're not familiar with the Azure service selected in the previous step, read the service documentation to understand the basics of the service.

App Service. A managed service for hosting web apps, mobile app back ends, RESTful APIs, or automated business processes.

Azure Kubernetes Service (AKS). A managed Kubernetes service for running containerized applications.

[Batch](#). A managed service for running large-scale parallel and high-performance

computing (HPC) applications

[Container Instances](#). The fastest and simplest way to run a container in Azure, without having to provision any virtual machines and without having to adopt a higher-level service.

[Functions](#). A managed FaaS service.

[Service Fabric](#). A distributed systems platform that can run in many environments, including Azure or on premises.

[Virtual machines](#). Deploy and manage VMs inside an Azure virtual network.

Understand the hosting models

Cloud services, including Azure services, generally fall into three categories: IaaS, PaaS, and FaaS. (There is also SaaS, software-as-a-service, which is out of scope for this article.) It's useful to understand the differences.

Infrastructure-as-a-Service (IaaS) lets you provision individual VMs along with the associated networking and storage components. Then you deploy whatever software applications you want onto those VMs. This model is the closest to a traditional on-premises environment, except that Microsoft manages the infrastructure. You still manage the individual VMs.

Platform-as-a-Service (PaaS) provides a managed hosting environment, where you can deploy your application without needing to manage VMs or networking resources. Azure App Service is a PaaS service.

Functions-as-a-Service (FaaS) goes even further in removing the need to worry about the hosting environment. In a FaaS model, you simply deploy your code and the service automatically runs it. Azure Functions are a FaaS service.

There is a spectrum from IaaS to pure PaaS. For example, Azure VMs can autoscale by using virtual machine scale sets. This automatic scaling capability isn't strictly PaaS, but is a type of management feature found in PaaS services.

In general, there is a tradeoff between control and ease of management. IaaS gives the most control, flexibility, and portability, but you have to provision, configure and manage the VMs and network components you create. FaaS services automatically manage nearly all aspects of running an application. PaaS services fall somewhere in between.

Consider limits and cost

- Service limits
- Cost
- SLA
- Regional availability

If on this analysis, you may find that the initial candidate isn't suitable for your particular application or workload. In that case, expand your analysis to include other Azure services.

Following tables contain additional comparison points, which may be useful when choosing.

Scaling model

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service
Application composition	Agnostic	Applications, containers	Services, guest executables, containers	Functions	Containers
Scalability	Agnostic	Multiple apps per instance via app service plans	Multiple services per VM	Serverless ¹	Multiple containers per node
Minimum number of nodes	1 ²	1	5 ³	Serverless ¹	3 ³
State management	Stateless or Stateful	Stateless	Stateless or stateful	Stateless	Stateless or Stateful

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service
Deployment model	Agnostic	Built in	Agnostic	Not applicable	Agnostic
Can be deployed to dedicated VMs?	Supported	Supported ⁵	Supported	Supported ⁵	Supported
Grid connectivity	Supported	Supported ⁶	Supported	Supported ⁷	Supported

5

If using Consumption plan. If using App Service plan, functions run on the VMs allocated for your App Service plan. See [Choose the correct service plan for Azure Functions](#).

Higher SLA with two or more instances.

Recommended for production environments.

Can scale down to zero after job completes.

Requires App Service Environment (ASE).

Use [Azure App Service Hybrid Connections](#).

Requires App Service plan.

vOps

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service
Deployment model	Agnostic	IIS Express, others ¹	Local node cluster	Visual Studio or Azure Functions CLI	Minikube, others
Programming model	Agnostic	Web and API applications, WebJobs for background tasks	Guest executable, Service model, Actor model, Containers	Functions with triggers	Agnostic

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service
Application lifecycle	No built-in support	Deployment slots	Rolling upgrade (per service)	Deployment slots	Rolling update

5

Options include IIS Express for ASP.NET or node.js (iisnode); PHP web server; Azure Toolkit for IntelliJ, Azure Toolkit for Eclipse. App Service also supports remote debugging of deployed web app.

See [Resource Manager providers, regions, API versions and schemas](#).

Availability

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service	Cloud Managed
Auto-scaling	Virtual machine scale sets	Built-in service	Virtual machine scale sets	Built-in service	Pod auto-scaling ¹ , cluster auto-scaling ²	Managed scaling
Load balancer	Azure Load Balancer	Integrated	Azure Load Balancer	Integrated	Azure Load Balancer or Application Gateway	Managed
Quota limit ³	Platform image: 1000 nodes per scale set, Custom image: 600 nodes per scale set	20 instances, 100 with App Service Environment	100 nodes per scale set	200 instances per Function app	100 nodes per cluster (default limit)	200 nodes per cluster (default limit)

5

See [Autoscale pods](#).

See [Automatically scale a cluster to meet application demands on Azure Kubernetes Service \(AKS\)](#).

See [Azure subscription and service limits, quotas, and constraints](#).

Availability

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service	Container Instances
Availability	SLA for Virtual Machines	SLA for App Service	SLA for Service Fabric	SLA for Functions	SLA for AKS	SLA for Container Instances
High availability	Traffic manager	Traffic manager	Traffic manager, Multi-Region Cluster	Azure Front Door	Traffic manager	Not supported

For guided learning on Service Guarantees, review [Core Cloud Services - Azure Architecture and service guarantees](#).

Other criteria

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service
Configuration	Configured in VM	Supported	Supported	Supported	Ingress controller
Cost	Windows, Linux	App Service pricing	Service Fabric pricing	Azure Functions pricing	AKS pricing
Architecture patterns	N-Tier, Big compute (HPC)	Web-Queue-Worker	Microservices, Event-driven architecture	Microservices, Event-driven architecture	Microservices, Event-driven architecture

Criteria	Virtual Machines	App Service	Service Fabric	Azure Functions	Azure Kubernetes Service

Output from this flowchart is a **starting point** for consideration. Next, perform a detailed evaluation of the service to see if it meets your needs.

Understand the basic features

If you're not familiar with the Azure service selected in the previous step, read one of the following overview articles:

- [App Service](#)
- [Azure Kubernetes Service](#)
- [Batch](#)
- [Container Instances](#)
- [Functions](#)
- [Service Fabric](#)
- [Virtual machines](#)

Consider limits and cost

To perform a more detailed evaluation, looking at the following aspects of the service:

- [Service limits](#)
- [Cost](#)
- [SLA](#)
- [Regional availability](#)
- [Compute comparison tables](#)

Based on this analysis, you may find that the initial candidate isn't suitable for your particular application or workload. In that case, expand your analysis to include other compute services.

Next steps

[Core Cloud Services - Azure compute options](#). This Microsoft Learn module explores how compute services can solve common business needs.

Is this page helpful?

Yes No

